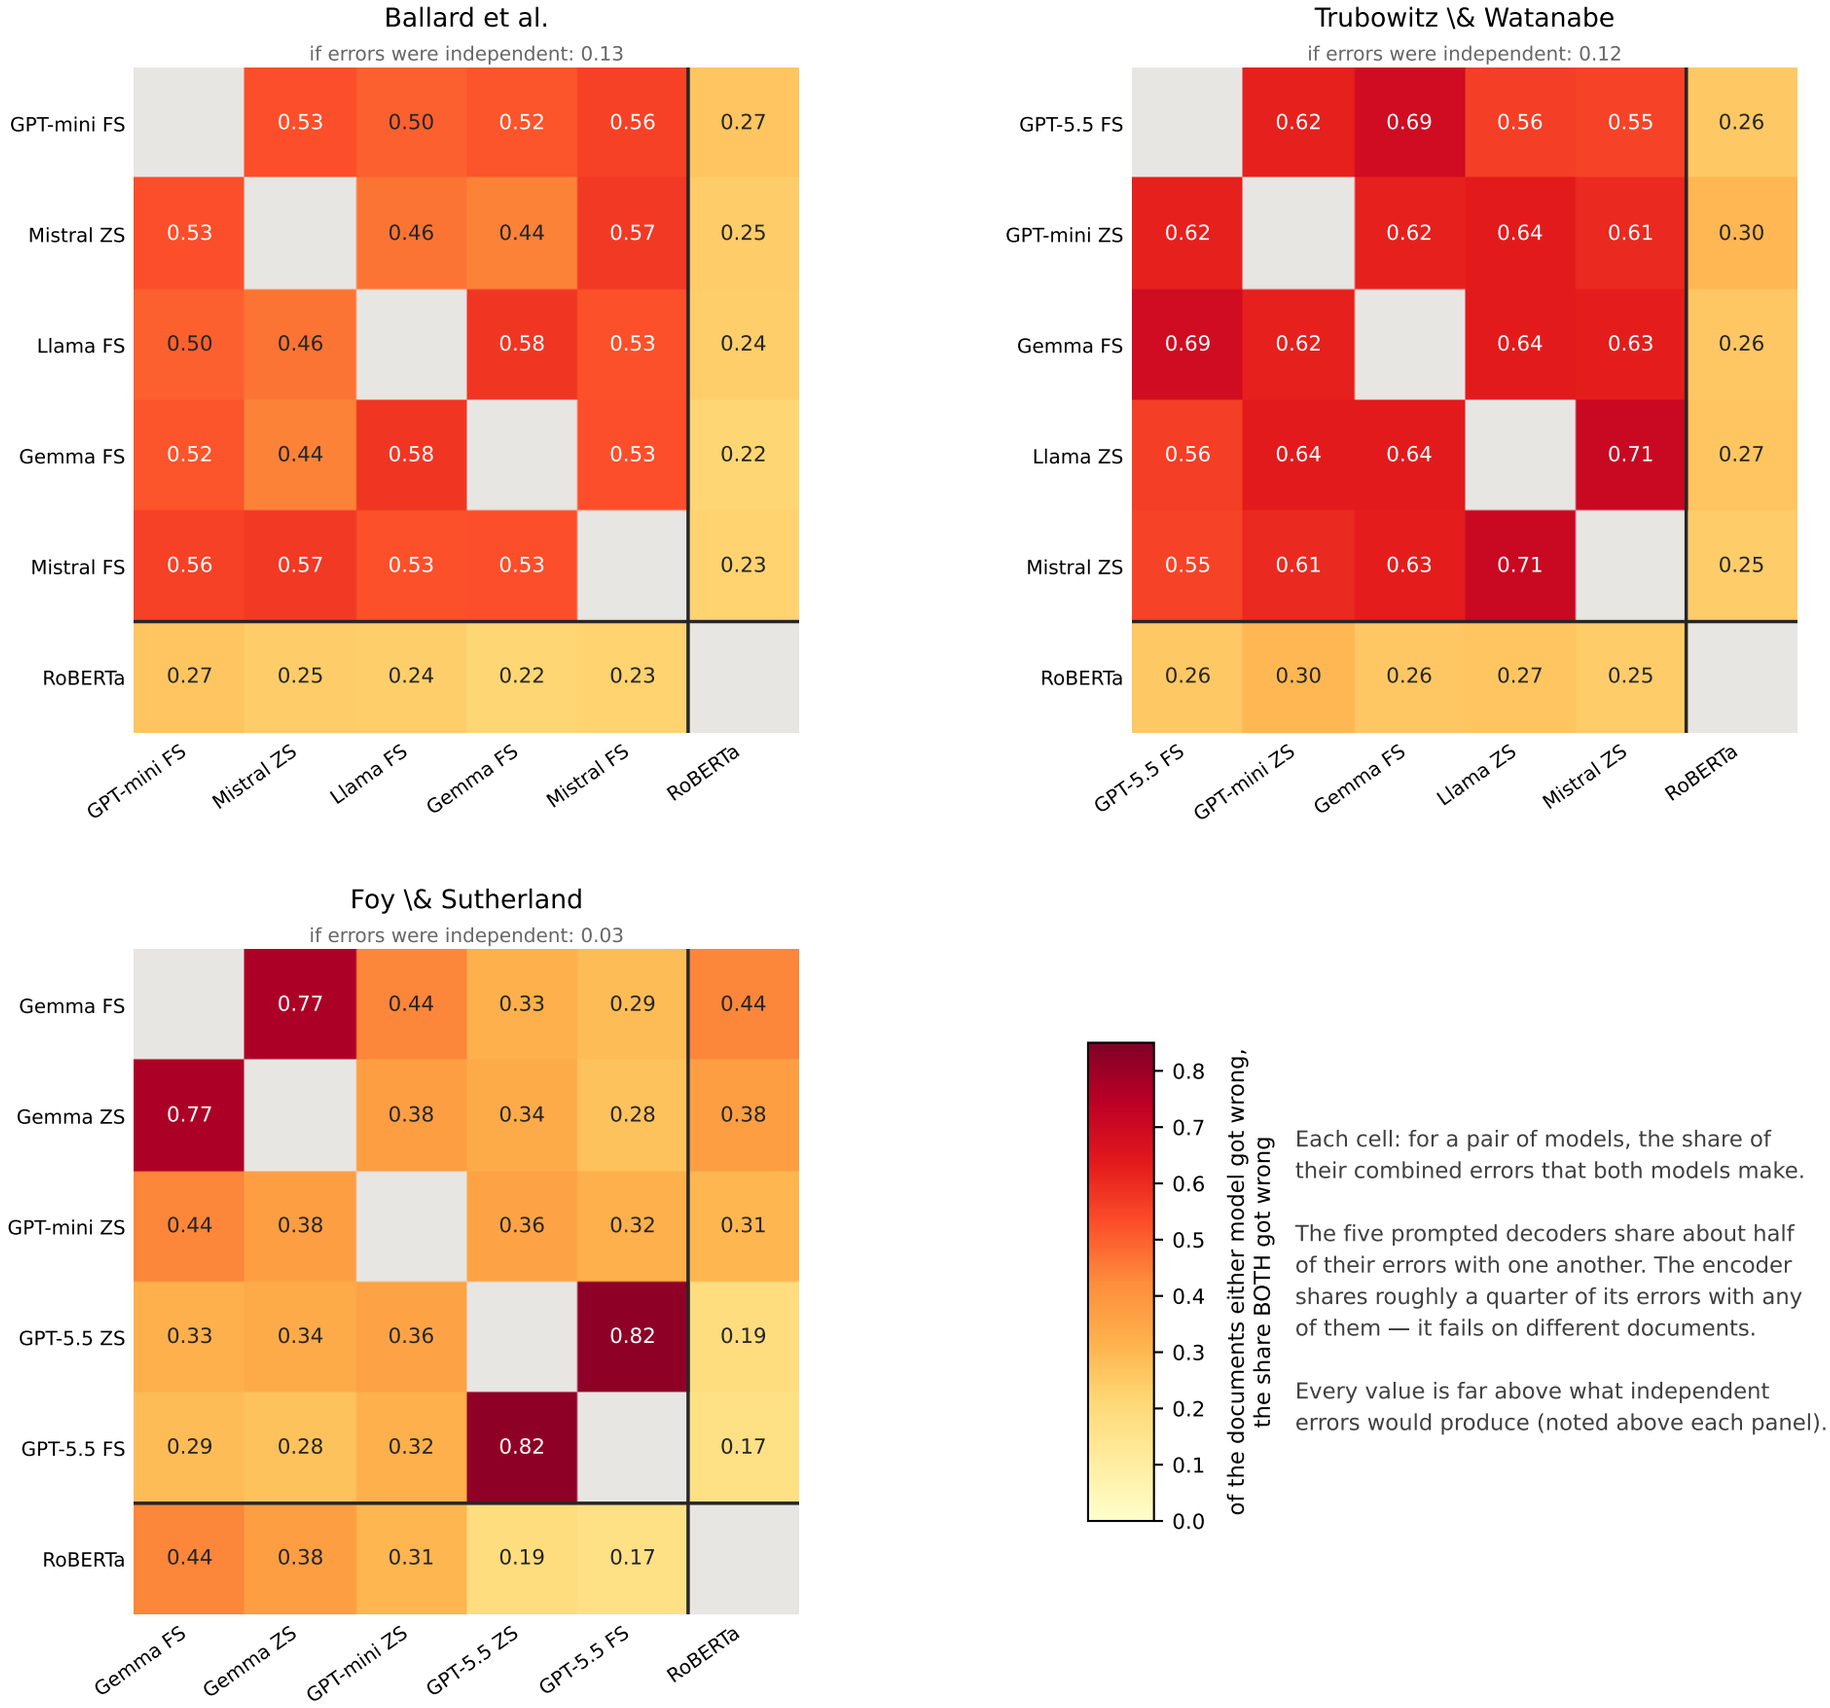
Ballard et al.
if errors were independent: 0.13
GPT-mini FS 0.53 0.50 0.52 0.56 0.27
Mistral ZS 0.53 0.46 0.44 0.57 0.25
Llama FS 0.50 0.46 0.58 0.53 0.24
Gemma FS 0.52 0.44 0.58 0.53 0.22
Mistral FS 0.56 0.57 0.53 0.53 0.23
RoBERTa 0.27 0.25 0.24 0.22 0.23
GPT-mini FS
Mistral ZS
Llama FS
Gemma FS
Mistral FS
RoBERTa
Trubowitz \& Watanabe
if errors were independent: 0.12
GPT-5.5 FS 0.62 0.69 0.56 0.55 0.26
GPT-mini ZS 0.62 0.62 0.64 0.61 0.30
Gemma FS 0.69 0.62 0.64 0.63 0.26
Llama ZS 0.56 0.64 0.64 0.71 0.27
Mistral ZS 0.55 0.61 0.63 0.71 0.25
RoBERTa 0.26 0.30 0.26 0.27 0.25
GPT-5.5 FS
GPT-mini ZS
Gemma FS
Llama ZS
Mistral ZS
RoBERTa
Foy \& Sutherland
if errors were independent: 0.03
Gemma FS 0.77 0.44 0.33 0.29 0.44
Gemma ZS 0.77 0.38 0.34 0.28 0.38
GPT-mini ZS 0.44 0.38 0.36 0.32 0.31
GPT-5.5 ZS 0.33 0.34 0.36 0.82 0.19
GPT-5.5 FS 0.29 0.28 0.32 0.82 0.17
RoBERTa 0.44 0.38 0.31 0.19 0.17
Gemma FS
Gemma ZS
GPT-mini ZS
GPT-5.5 ZS
GPT-5.5 FS
RoBERTa
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
of the documents either model got wrong, the share BOTH got wrong
Each cell: for a pair of models, the share of their combined errors that both models make.
The five prompted decoders share about half of their errors with one another. The encoder shares roughly a quarter of its errors with any of them — it fails on different documents.
Every value is far above what independent errors would produce (noted above each panel).